

Assignment -2

DATA VISUALISATION AND PRE-PROCESSING

Assignment Date	28 September 2022
Student Name	Anitha A
Student Roll Number	612419104004
Maximum Marks	2 Marks

Question-1:

Download the Dataset

Solution:

[illegible]

Question-2:

Loading dataset

Solution:

```
df = pd.read_csv("Churn_Modelling.csv")
```

df

```
In [2]: df = pd.read_csv("Churn_Modelling.csv")
```

```
In [3]: df
```

```
Out[3]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	Spain	Female	41	1	81807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0
--	--	--	--	--	--	--	--	--	--	--	--	--	--	--
9995	9996	15606329	Obijaku	771	France	Male	39	5	0.00	2	1	0	96270.64	0
9996	9997	15609932	Johnstone	516	France	Male	35	10	57369.61	1	1	1	101609.77	0
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	1	42085.58	1
9998	9999	15682955	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	0	92888.52	1
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1	0	38190.78	0

10000 rows x 14 columns

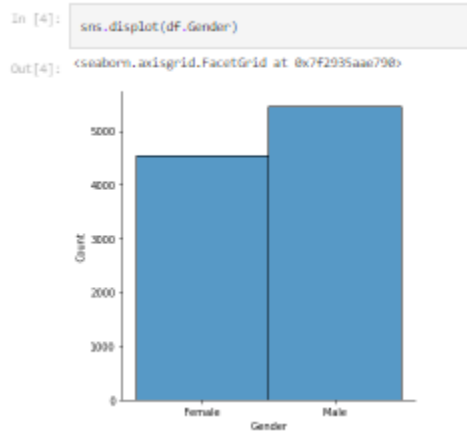
Question-3:

1. Visualizations

a) Univariate Analysis

Solution:

```
sns.displot(df.Gender)
```




b) Bi - Variate Analysis

Solution:

```
df.plot.line()
```

```
In [5]: df.plot.line()
Out[5]: <matplotlib.axes._subplots.AxesSubplot at 0x7F2932d2dF58>
```




C)Multi - Variate Analysis

Solution:

```
sns.lmplot("Tenure", "NumOfProducts", df, hue="NumOfProducts", fit_reg=False);
```

```
In [6]: sns.lmplot("Tenure", "NumOfProducts", df, hue="NumOfProducts", fit_reg=False);
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y, data. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.



Question-4:

Perform descriptive statistics on the dataset.

Solution:

```
df.describe()
```

```
In [7]: df.describe()
```

```
Out[7]:
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.526600	38.921800	5.012800	76485.889288	1.530200	0.70550	0.515100	100090.239881	0.201700
std	2886.89566	7.193619e+04	96.653299	10.487806	2.892174	62397.405032	0.581654	0.45584	0.499797	57510.492818	0.402769
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000000	11.580000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.000000	51002.110000	0.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97196.540000	1.000000	1.00000	1.000000	100191.915000	0.000000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000	1.000000	149388.247500	0.000000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.00000	1.000000	199992.480000	1.000000

Question-5:

Handle the missing values

Solution:

```
data = pd.read_csv("Churn_Modelling.csv")
pd.isnull(data["Gender"])
```

```
In [8]: data = pd.read_csv("Churn_Modelling.csv")
pd.isnull(data["Gender"])
```

```
Out[8]:
```

```
0      False
1      False
2      False
3      False
4      False
...
9995   False
9996   False
9997   False
9998   False
9999   False
Name: Gender, Length: 10000, dtype: bool
```

Question-6:

Find the outliers and replace the outliers.


Solution:

```
sns.boxplot(df['Age'])
```

```
In [9]: sns.boxplot(df['Age'])
```

```
FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be "data", and passing other arguments without an explicit keyword will result in an error or misinterpretation.
```

```
Out[9]:
```



Solution:

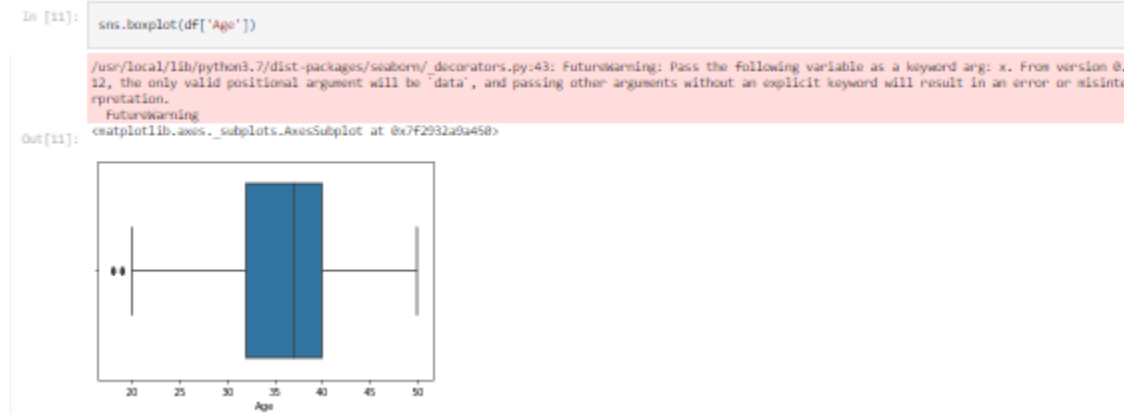
```
df['Age']=np.where(df['Age']>50,40,df['Age'])
df['Age']
```

```
In [10]: df['Age']=np.where(df['Age']>50,40,df['Age'])
df['Age']

Out[10]: 0    42
1    41
2    42
3    39
4    43
..
9995  39
9996  35
9997  36
9998  42
9999  28
Name: Age, Length: 10000, dtype: int64
```

Solution:

```
sns.boxplot(df['Age'])
```



Solution:

```
df['Age']=np.where(df['Age']<20,35,df['Age'])
df['Age']
```

```
In [12]: df['Age']=np.where(df['Age']<20,35,df['Age'])
df['Age']

Out[12]: 0    42
1    41
2    42
3    39
4    43
..
9995  39
9996  35
9997  36
9998  42
9999  28
Name: Age, Length: 10000, dtype: int64
```

Question-7:

Check for Categorical columns and perform encoding.

Solution:

```
pd.get_dummies(df, columns=["Gender", "Age"], prefix=["Age", "Gender"]).head
```

```
In [13]: pd.get_dummies(df, columns=["Gender", "Age"], prefix=["Age", "Gender"]).head()
```

```
Out[13]:
```

	nce	NumOfProducts	HasCrCard	IsActiveMember	..	Gender_41	Gender_42	Gender_43	Gender_44	Gender_45	Gender_46	Gender_47	Gender_48	Gender_49	Gender_50
0.00	1	1	1	1	..	0	1	0	0	0	0	0	0	0	0
7.86	1	0	1	1	..	1	0	0	0	0	0	0	0	0	0
0.80	3	1	0	1	..	0	1	0	0	0	0	0	0	0	0
0.00	2	0	0	1	..	0	0	0	0	0	0	0	0	0	0
0.82	1	1	1	1	..	0	0	1	0	0	0	0	0	0	0

Question-8:

Split the data into dependent and independent variables.

a)split the data into independent variables.

Solution:

```
X = df.iloc[:, :-1].value
print(X)
```

```
In [14]: X = df.iloc[:, :-1].values
print(X)

[[1 15634682 'Hargrave' ... 1 1 101348.88]
 [2 15647311 'Hill' ... 0 1 112542.58]
 [3 15619384 'Onio' ... 1 0 113931.57]
 ...
 [9998 15584532 'Liu' ... 0 1 42885.58]
 [9999 15682355 'Sabbatini' ... 1 0 92888.52]
 [10000 15628319 'Walker' ... 1 0 38190.78]]
```

b)split the data into dependent variables

Solution:

```
Y = df.iloc[:, -1].values
print(Y)
```

```
In [15]: Y = df.iloc[:, -1].values
print(Y)

[1 0 1 ... 1 1 0]
```

Question-9:

Scale the independent variables

Solution:

```
import pandas as pd
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
df[["CustomerId"]] = scaler.fit_transform(df[["CustomerId"]])
```

```
In [16]: import pandas as pd
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
df[["CustomerId"]] = scaler.fit_transform(df[["CustomerId"]])

In [17]: print(df)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
0	1	0.275616	Hargrave	619	France	Female	42	
1	2	0.326454	Hill	688	Spain	Female	41	
2	3	0.214421	Onio	582	France	Female	42	
3	4	0.542636	Boni	699	France	Female	39	
4	5	0.688778	Mitchell	850	Spain	Female	43	
...
9995	9996	0.162119	Obijaku	771	France	Male	39	
9996	9997	0.816765	Johnstone	516	France	Male	35	
9997	9998	0.875327	Liu	789	France	Female	36	
9998	9999	0.466637	Sabbatini	772	Germany	Male	42	
9999	10000	0.250483	Walker	792	France	Female	28	

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83887.86	1	0	1	
2	8	159668.88	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	
...
9995	5	0.00	2	1	0	
9996	10	57369.61	1	1	1	
9997	7	0.00	1	0	1	
9998	3	75075.31	2	1	0	
9999	4	130142.79	1	1	0	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0
...
9995	96270.64	0
9996	101699.77	0
9997	42085.58	1
9998	92888.52	1
9999	38190.78	0

[10000 rows x 14 columns]

Question-10:

Split the data into training and testing

Solution:

```
from sklearn.model_selection import train_test_split
train_size=0.8
X = df.drop(columns = ['Tenure']).copy()
y = df['Tenure']
X_train, X_rem, y_train, y_rem = train_test_split(X,y, train_size=0.8)
test_size = 0.5
X_valid, X_test, y_valid, y_test = train_test_split(X_rem,y_rem, test_size=0.5)
print(X_train.shape), print(y_train.shape)
print(X_valid.shape), print(y_valid.shape)
print(X_test.shape), print(y_test.shape)
```

```
In [18]: from sklearn.model_selection import train_test_split
train_size=0.8
X = df.drop(columns = ['tenure']).copy()
y = df['tenure']
X_train, X_rem, y_train, y_rem = train_test_split(X,y, train_size=0.8)
test_size = 0.5
X_valid, X_test, y_valid, y_test = train_test_split(X_rem,y_rem, test_size=0.5)
print(X_train.shape), print(y_train.shape)
print(X_valid.shape), print(y_valid.shape)
print(X_test.shape), print(y_test.shape)

(8000, 13)
(8000,)
(1000, 13)
(1000,)
(1000, 13)
(1000,)
Out[18]: (None, None)
```